

JRA 2/3 Infrastructure Management Workshop

Seamless RI integration using declarative configuration management

Steffen Vogel

Institute for Automation of Complex Power Systems

RWTH Aachen University, Aachen, Germany

26th January 2022



© The ERIGrid 2.0 Consortium
EU H2020 Programme GA No. 870620

● The workshop will be recorded

Agenda

1. Problem Statement
 - Objectives for today
2. Infrastructure as Code
 - Declarative Configuration Management
3. Kubernetes
 - Containers
 - Interfaces
 - Operators & Reconciliation
4. RIasC
 - Provisioning
 - Virtual Networking
 - Network Emulation
 - Network Monitoring
 - Time-synchronization
 - Multi Tenancy
5. Operator Example
6. Feedback, Discussion & Questions



Duration:
1-2 hrs



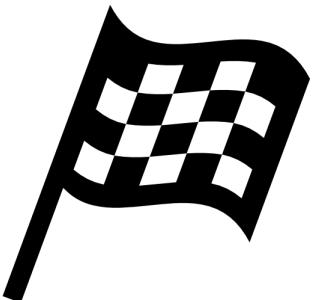
Problem & Challenges

- Infrastructure Management
 - Complexity
 - Reproducability
 - Drift
 - Resource Sharing
 - Remote Access & Remote Preparation



Objectives of this Workshop

- Present JRA2.2.5 development “*Accelerate time-to-experiment for Remote RI*”
- Introduce the basics of Kubernetes & declarative configuration management
 - The operator reconciliation loop
- Provide inputs and ideas for ERIGrid’s JRA3 platform
 - RIasC - Research Infrastructure as Code



Inspiration: Infrastructure as Code



“ **Infrastructure as code (IaC)** is the process of managing and provisioning computer data centers through machine-readable definition files, rather than physical hardware configuration or interactive configuration tools.

The infrastructure managed by this process comprises both physical equipment, as well as software tools. The definitions may be in a version control system. It can use either scripts or declarative definitions, rather than manual processes, but the term is more often used to promote declarative approaches.”

Source: [Wikipedia](#)



Research Infrastructure as Code (RIsaC)



“ **Research Infrastructure as code (RIsaC)** is the process of managing and provisioning *research experiments* through machine-readable definition files, rather than manual hardware configuration or interactive configuration tools.

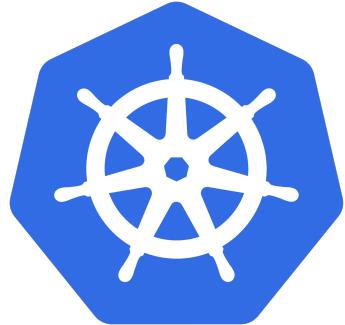
Experiments managed by this process comprises both physical equipment, as well as software tools. The definitions may be in a version control system. It can use either scripts or declarative definitions, rather than manual processes, but the term is more often used to promote declarative approaches.”



Existing tools for IaC / Configuration Management

- Puppet (2005)
- Chef (2009)
- SaltStack (2011)
- Ansible (2012)
- Terraform (2014)
- Kubernetes (2015)





Container Orchestration & more?

KUBERNETES



Kubernetes

- Container Orchestration System
 - Executes of containerized applications on a cluster of one or multiple machines
 - Uses declarative configuration for cluster state
- Open Sourced by Google in 2015
 - Now managed by Cloud Native Computing / Linux Foundation
 - ~4 releases per year
 - Written in Go programming language
- Highly extensible
 - Plugins / Operators for Storage, Networking, and much more
 - Custom Ressources



“Everything at Google runs in a Container”

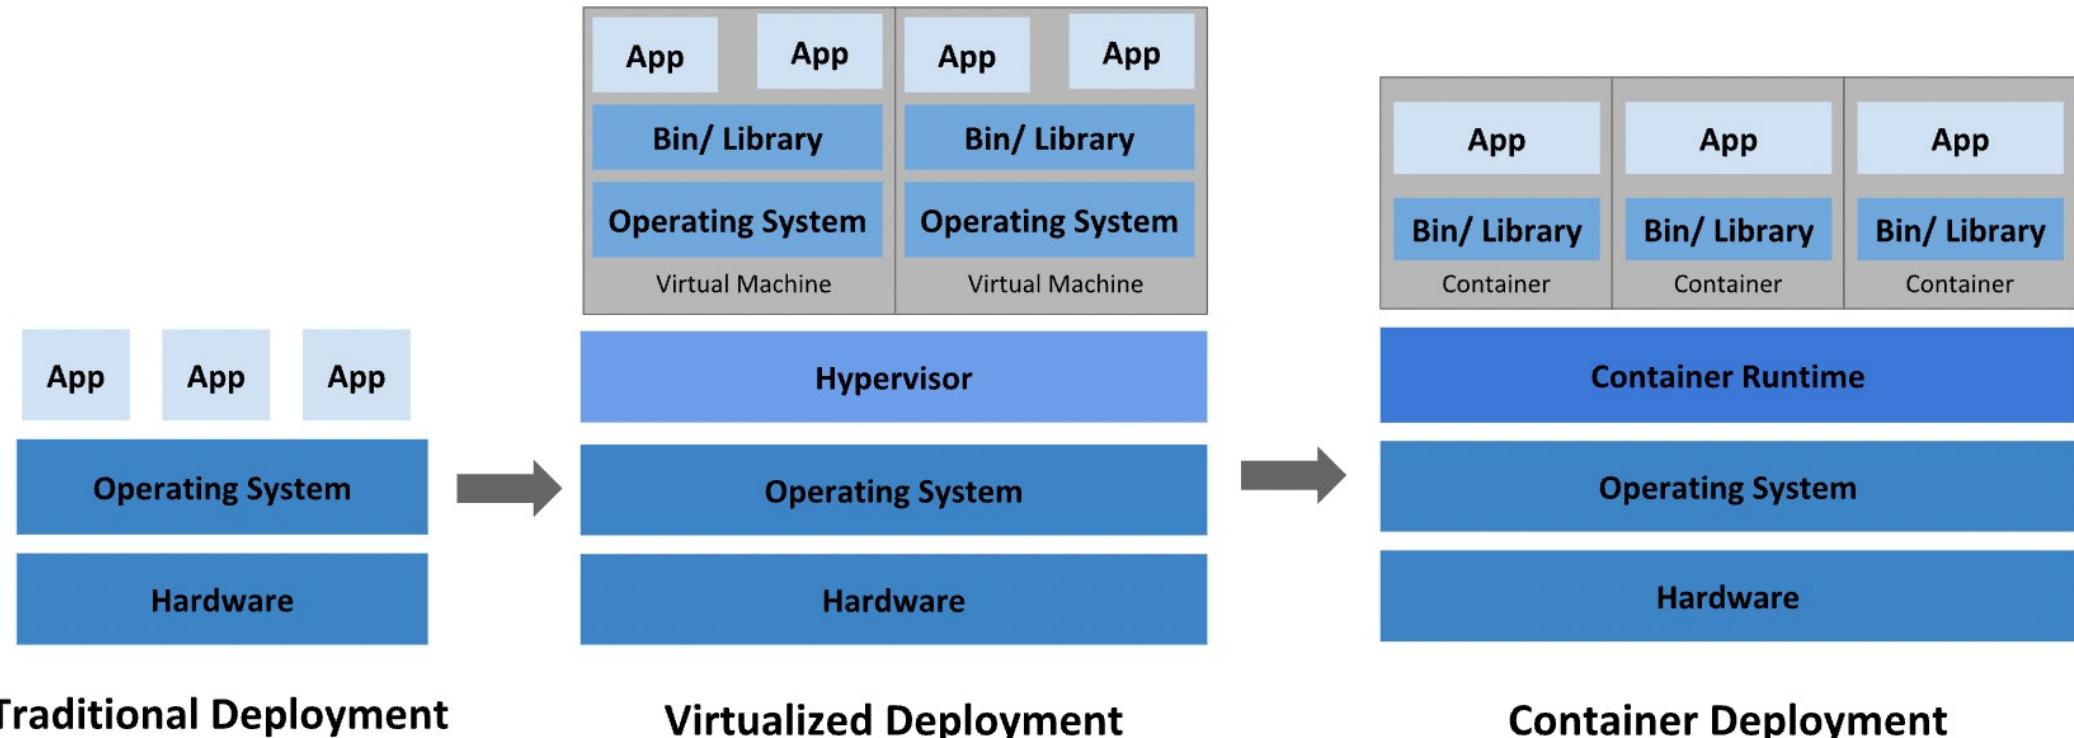
“We start over 2 billion containers per week”

(around 3000 per second)

Source: “[Joa Beda: Containers as Scale](#), GlueCon, May 22 2014”



Evolution of software deployment



Source: <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>



Containers & Images

- Packages software applications in an **exchangable** and **reproducible** format
- Large library of existing images for existing software
- Dockerfile's as recipes for building a Container image:

Dockerfile

```
FROM python:3.8
WORKDIR /
ADD requirements.txt main.py /
RUN pip3 install -r requirements.txt
CMD [ "python3" , "/main.py" ]
```

Test locally:

```
docker build -t erigrid/py_ctrl:v1 -f Dockerfile .
docker run erigrid/py_ctrl:v1
```

Publish or get to image registry:

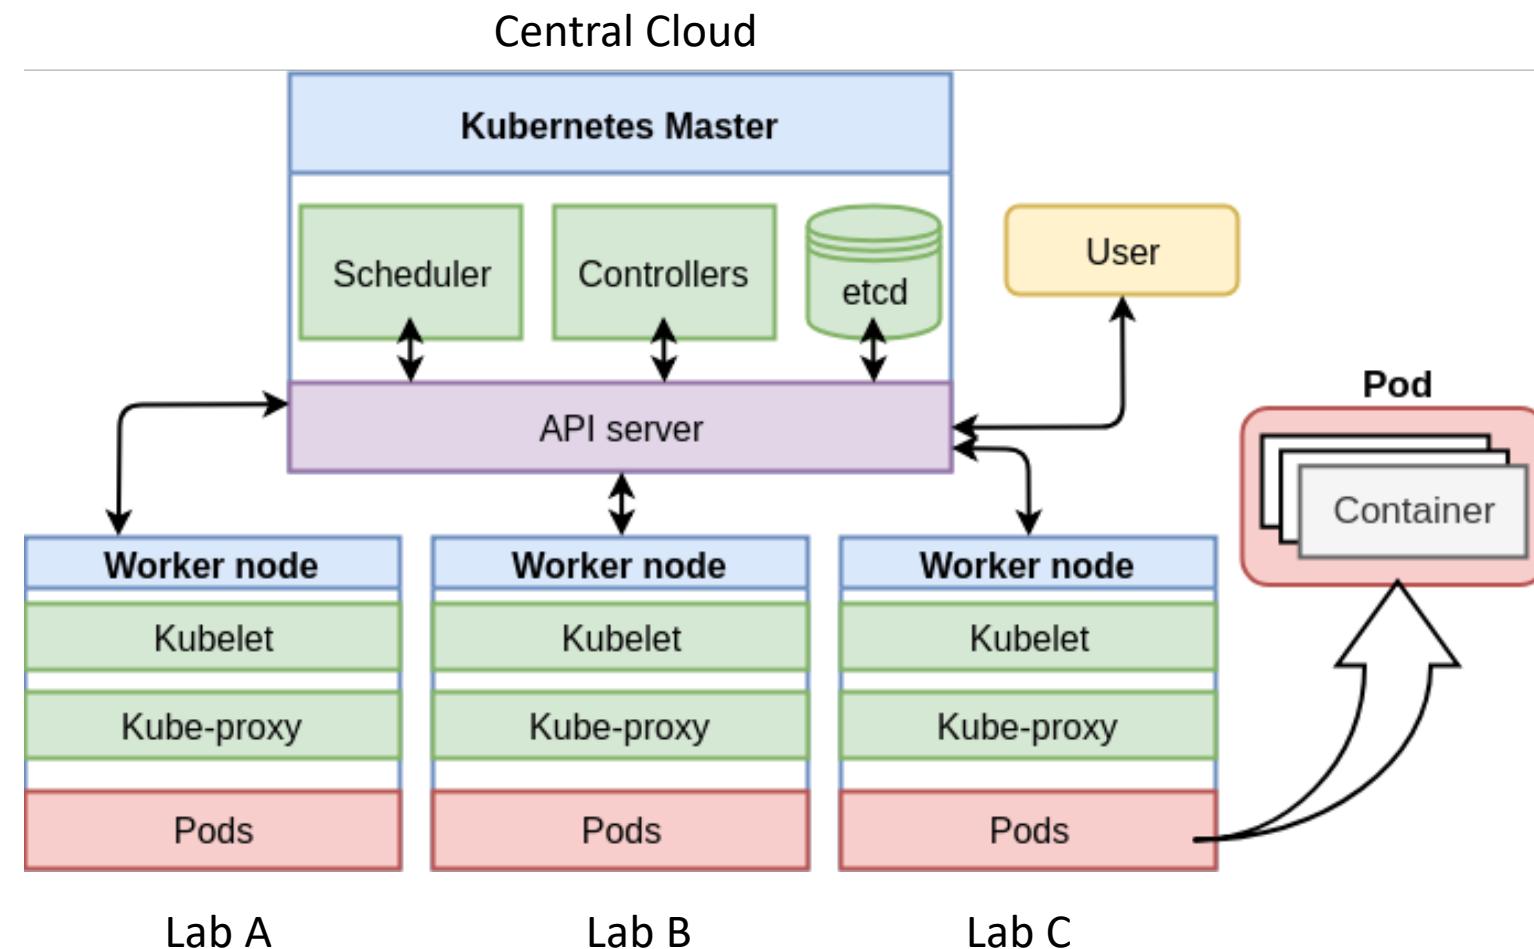
```
docker push erigrid/my_py_ctrl:v1
docker pull erigrid/my_py_ctrl:v1
```



ERIGrid Images: <https://hub.docker.com/orgs/erigrid/repositories>



Architecture

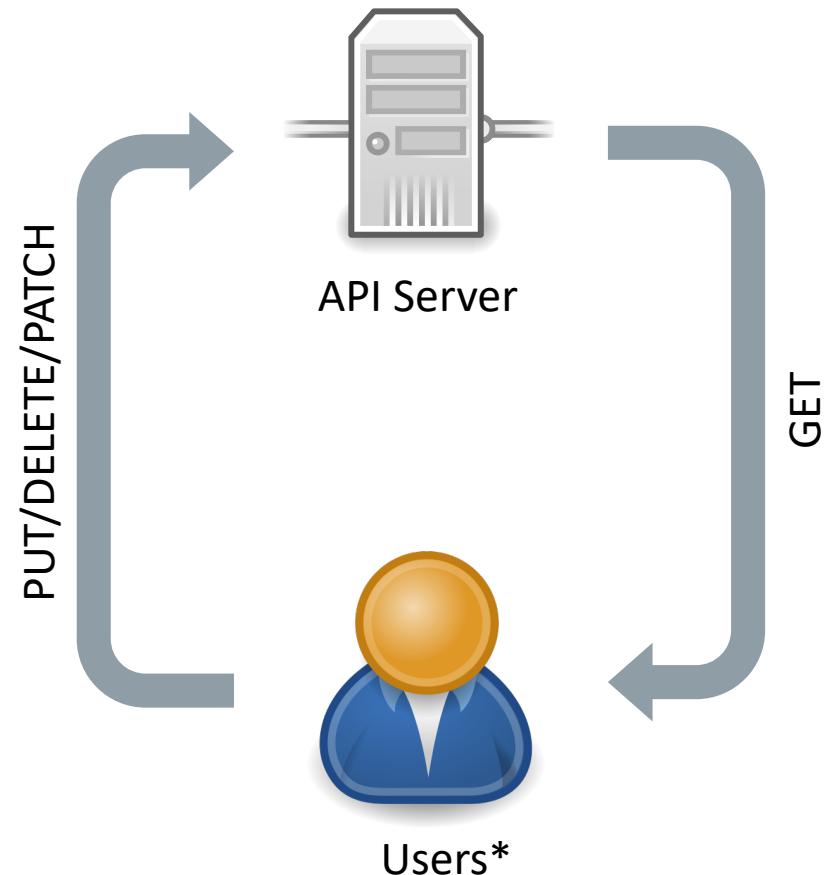


Source: "Olle Larsson, Running Databases in a Kubernetes Cluster, M.Sc. Thesis, 2019, Umea University"



Manifests

- Each resource in Kubernetes is described by a manifest
 - Represents the **desired** state
 - Many types of resources are available:
 - Workloads: Containers, Pods
 - Storage: Storage Volumes, Configuration Files, Secrets
 - Network: Load Balancers, Autoscaling ...
- They are stored by Kubernetes API server
- Usually formatted as JSON or YAML documents



Manifests: Basic Structure

```
---  
apiVersion: v1  
kind: Pod  
metadata:  
  name: py_ctrl  
  labels:  
    app: controller  
    project: eg2-ta-123  
annotations:  
  last-change: "2022-01-24 11:01:01"  
spec:  
  [ ... ]
```

YAML Document Separator
(when describing multiple resources in a single file)

Preamble
(same for all resources)

Resource Specification
(structure depends on **kind** field)



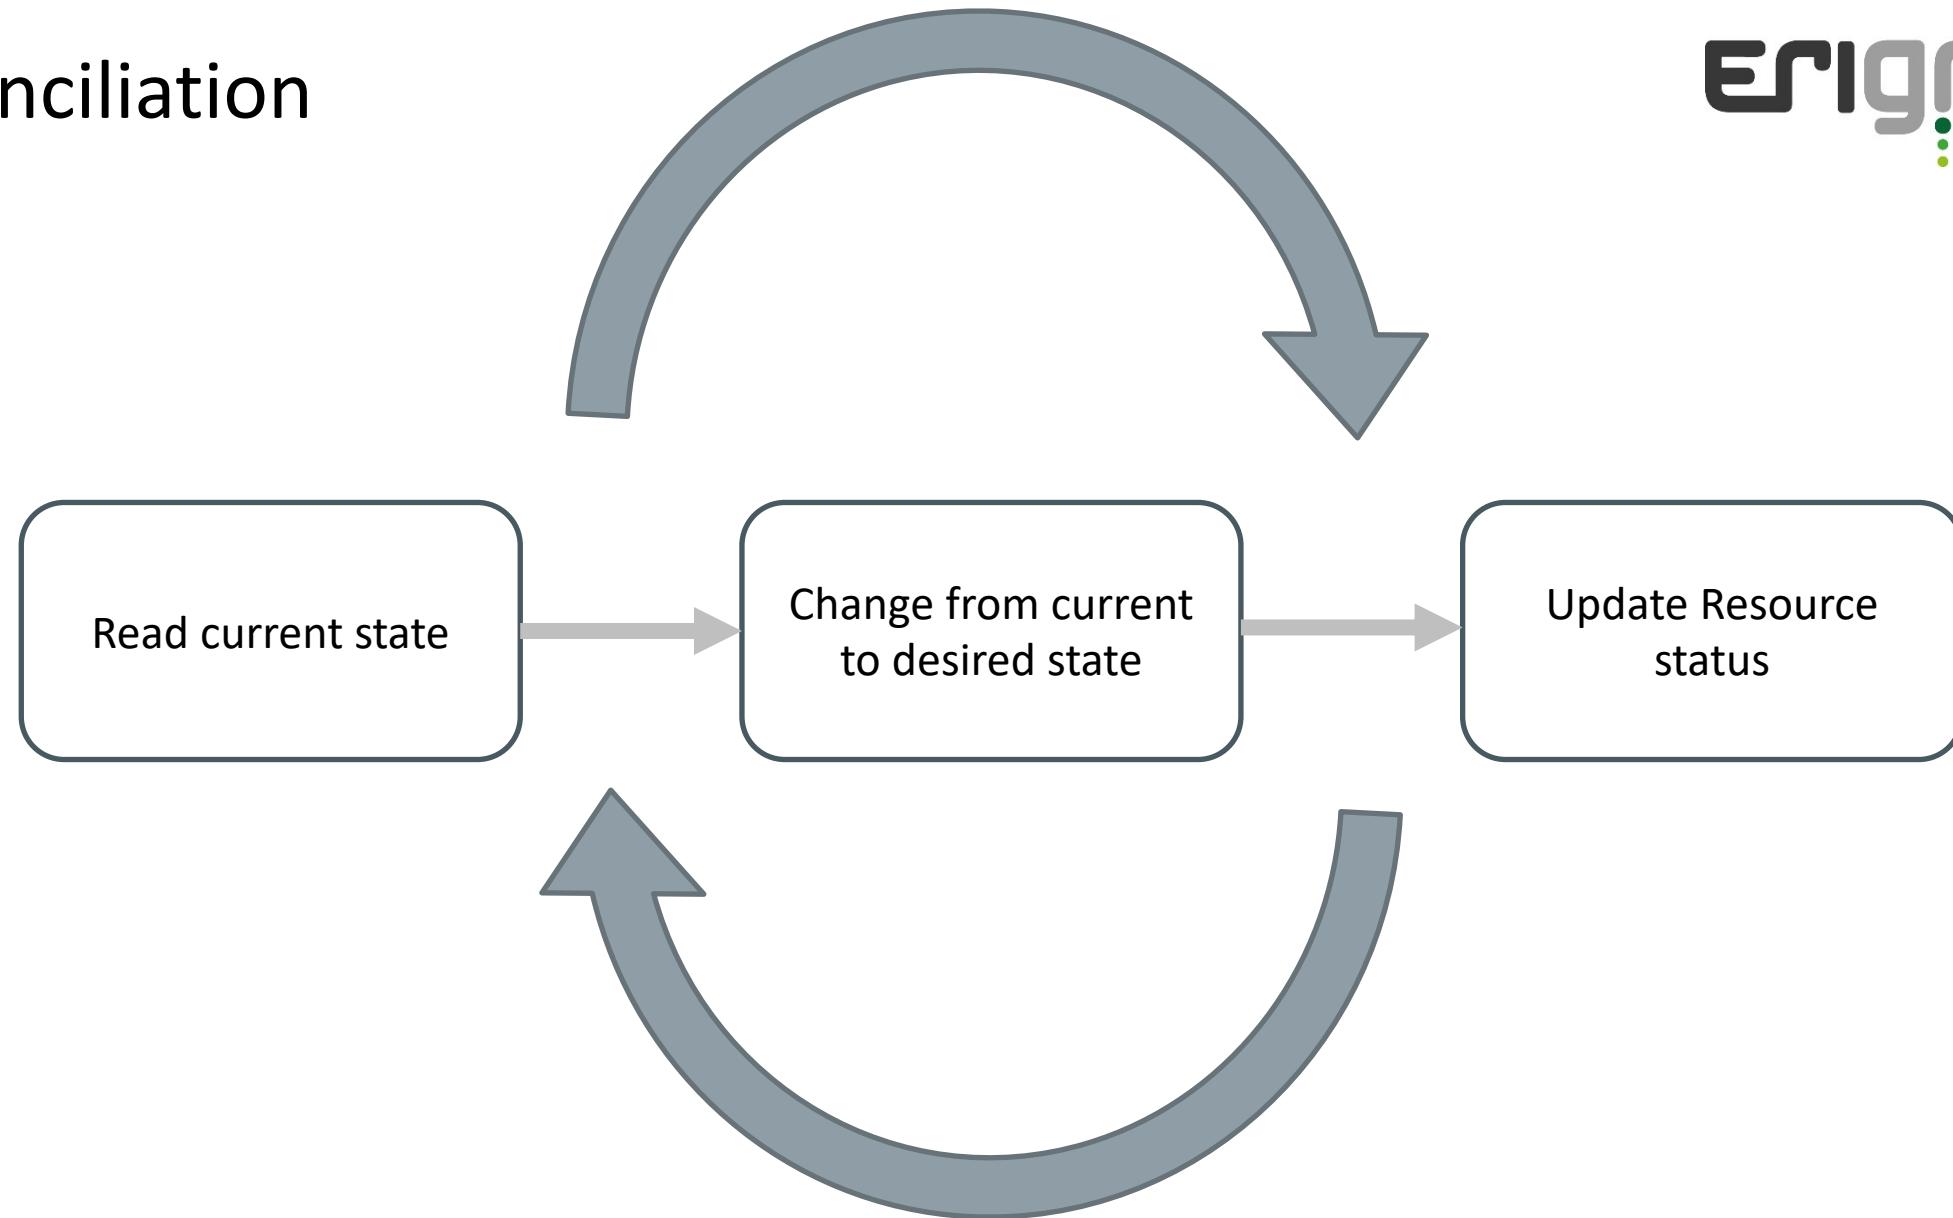
Manifests: Examples

```
---  
apiVersion: v1  
kind: Pod  
metadata:  
  name: py_ctrl  
  labels:  
    app: controller  
    project: eg2-ta-123  
  annotations:  
    last-change: "2022-01-24 11:01:01"  
spec:  
  nodeSelector:  
    institution: rwth  
  containers:  
  - name: controller  
    image: erigrid/py_ctrl
```

```
---  
apiVersion: v1  
kind: PersistentVolumeClaim  
metadata:  
  name: results  
spec:  
  accessModes:  
  - ReadWriteOnce  
  resources:  
    requests:  
      storage: 8Gi
```

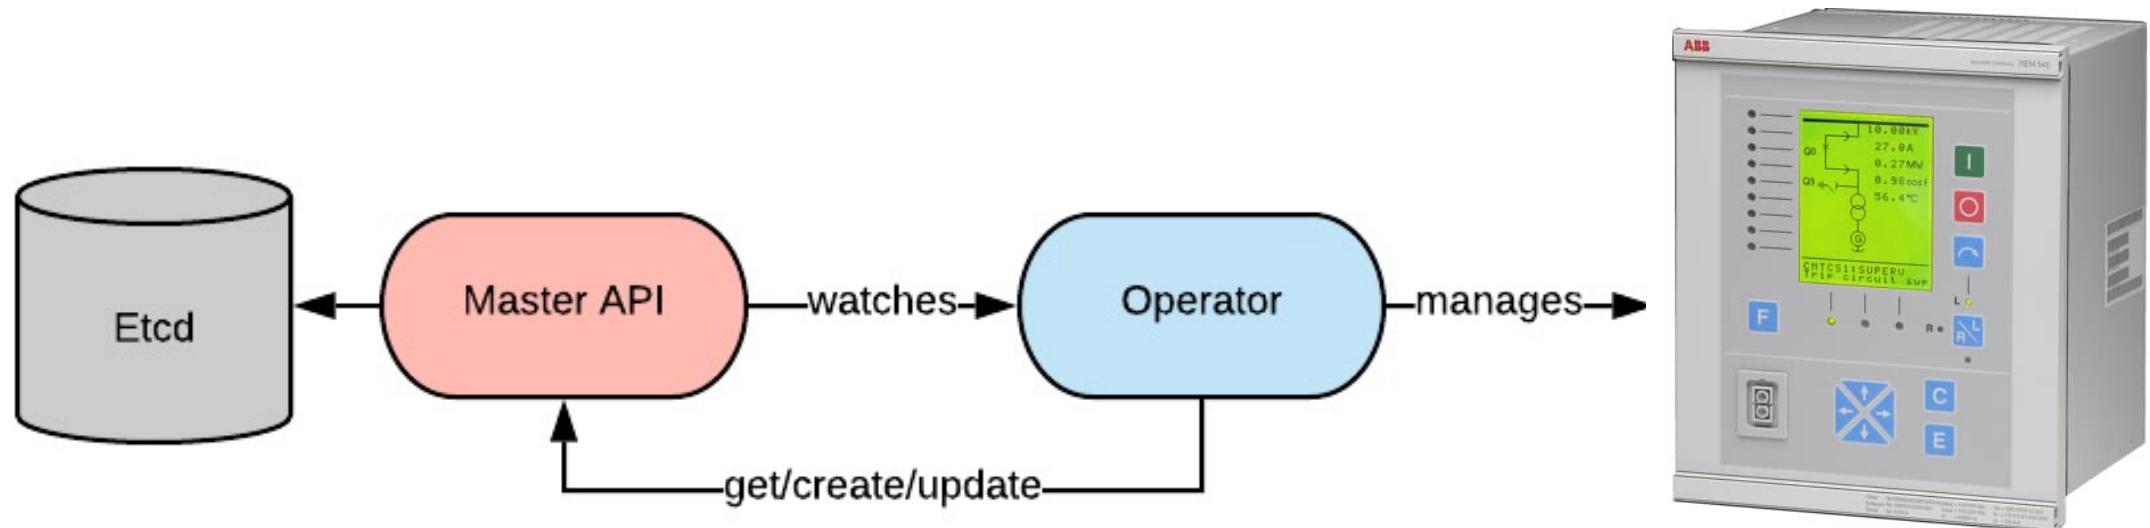


Reconciliation



Custom Resource Descriptions (CRD)

- We can define our own resources types
 - And logic (operators) which processes them



Source: <https://cloud.redhat.com/blog/kubernetes-operators-best-practices>



Interfaces

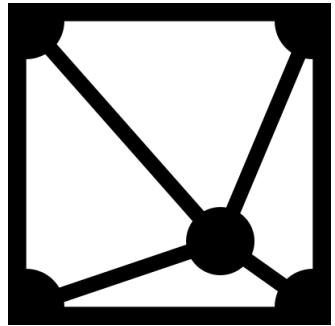
- **Web**
 - Kubernetes Dashboard
 - Rancher UI
- **CLI: Command Line**
 - kubectl
 - helm
 - curl
- **API: REST / Language Bindings**
- **IDE: Development Environments**
 - Lens
 - Visual Studio Code
- **Custom?**



Kubernetes REST API Specification:

<https://petstore.swagger.io/?url=https://raw.githubusercontent.com/kubernetes/kubernetes/master/api/openapi-spec/swagger.json#/>





Research Infrastructure as Code

RiasC
riasc.eu

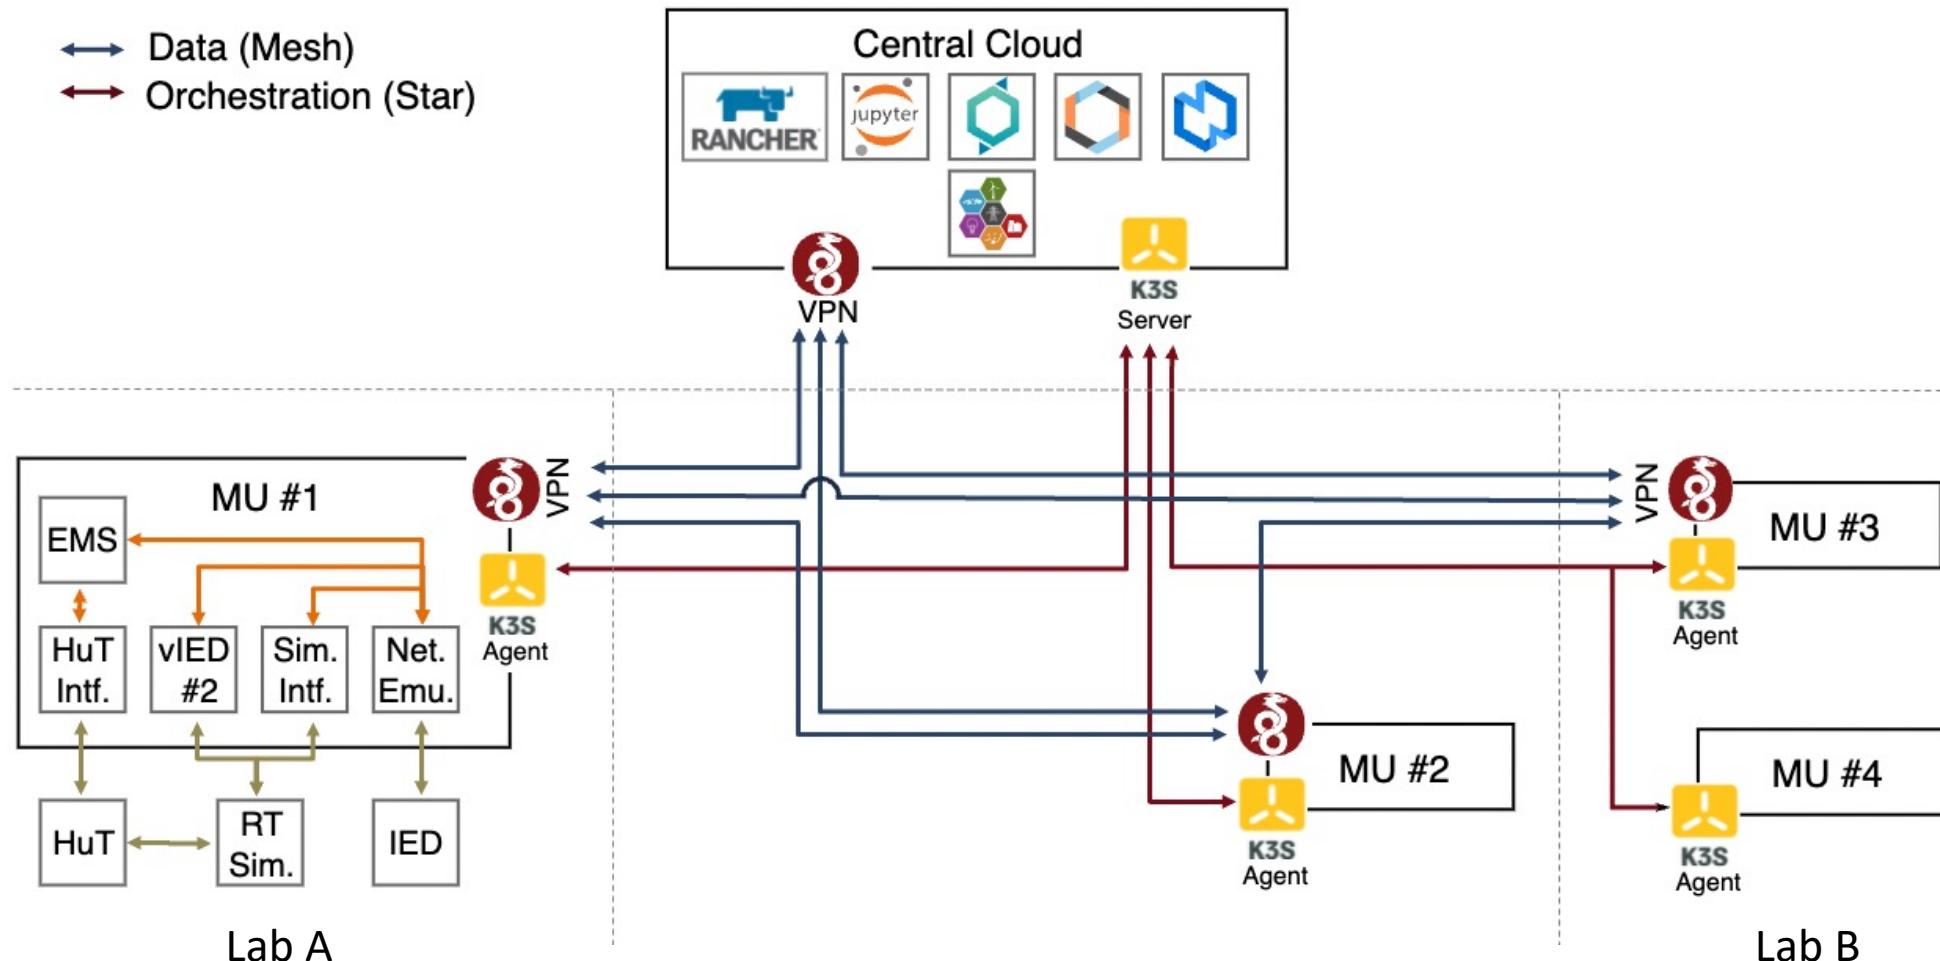


- RlasC is a set of tools to run research experiments on a (distributed) Kubernetes cluster
 - Opinionized selection of Kubernetes addons
 - Networking, Storage
 - Custom Resources
 - Operators



Architecture

↔ Data (Mesh)
↔ Orchestration (Star)



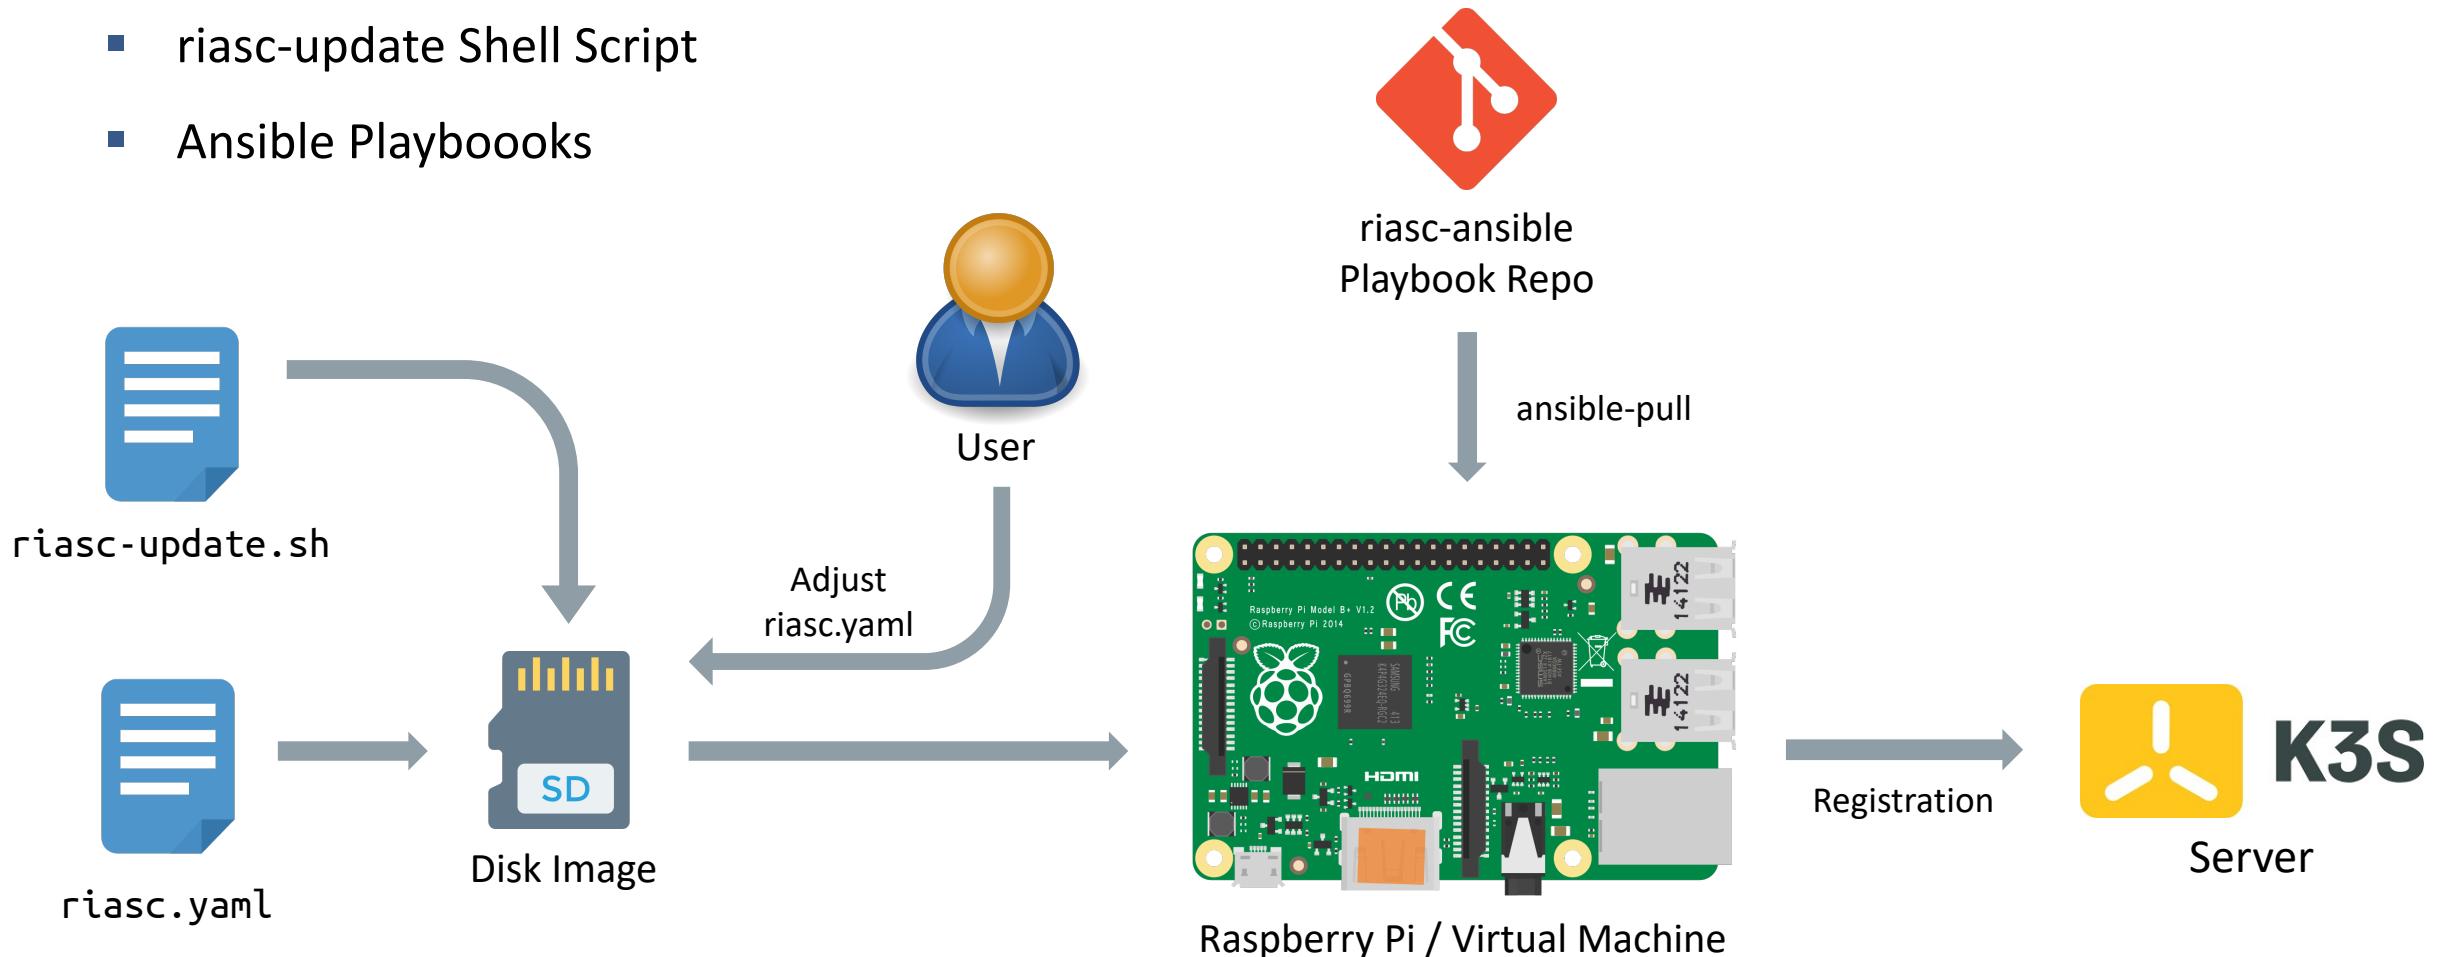
1. Provisioning of Cluster Nodes

- Streamlined process for adding new nodes into the cluster
 - Usually takes 15min to setup a new node
- Requirements:
 - Raspberry Pi or Debian/Redhat-based distro
 - Internet connectivity
 - Hostname and Token
- Recipe
 - 1. Download image from [riasc.eu](#)
 - 2. Copy to SDcard
 - 3. Adjust `riasc.yaml` configuration file on /boot partition
 - 4. Power-on device and wait...



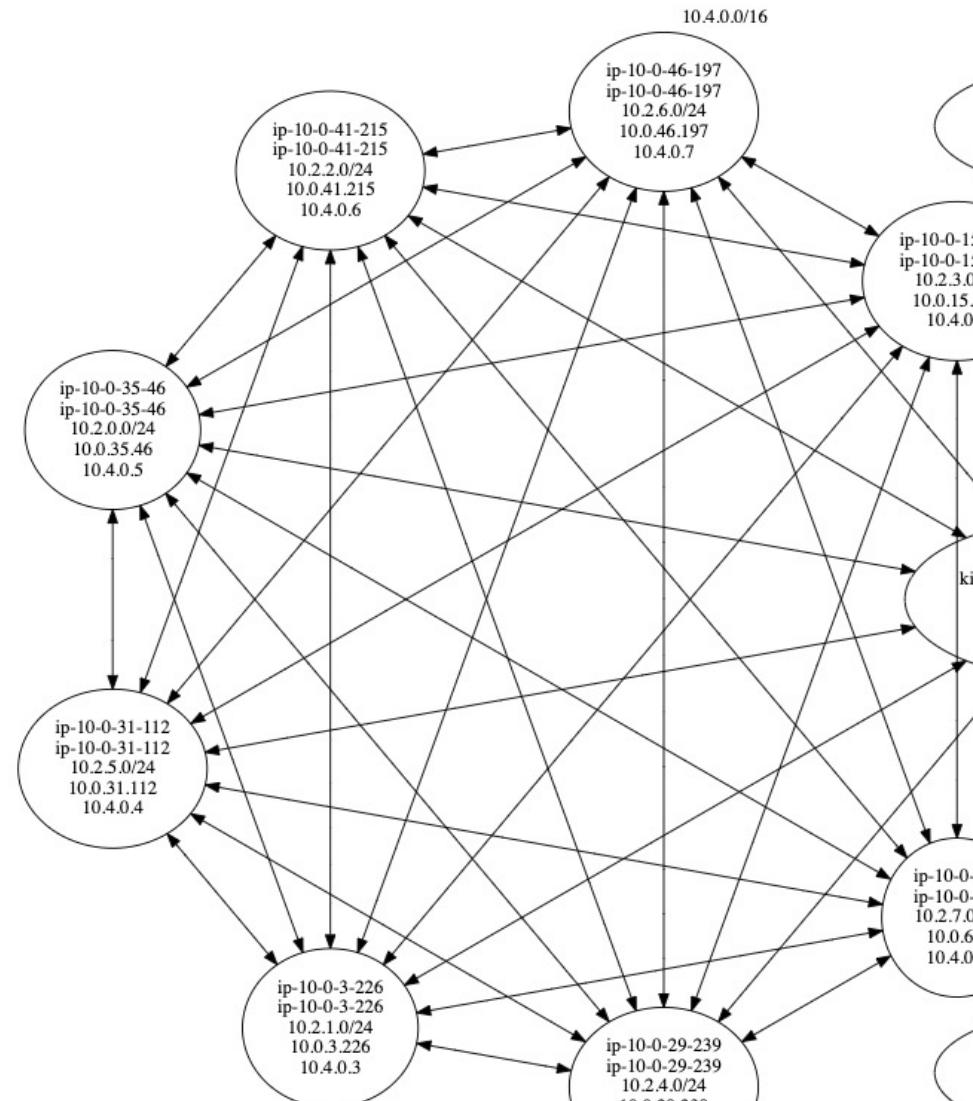
1. Provisioning /2

- Automated SDcard image builder
- riasc-update Shell Script
- Ansible Playbooks



2. Virtual Networking

- All Containers* in the cluster are part of the same cluster internal network
 - Each container* gets a dedicated IP address assigned
- Usually not a problem for centralized clusters
 - However, requires some addons for distributed clusters
- Kubernetes addons: [Kilo](#) & Wireguard
- Automatic NAT traversal



2. Network Emulation

- Allows the user to declaratively define network conditions between Pods in the cluster
- Uses Linux Traffic Control subsystem (netem Queuing Discipline)
- Live update of filters without restart of Pods
- Multiple “controllers” available
 - Builtin: simple tc-netem settings
 - Flexe: VTT’s network emulation tool
 - Supports segment sequencing to simulate alternating network conditions



2. Network Emulation /2

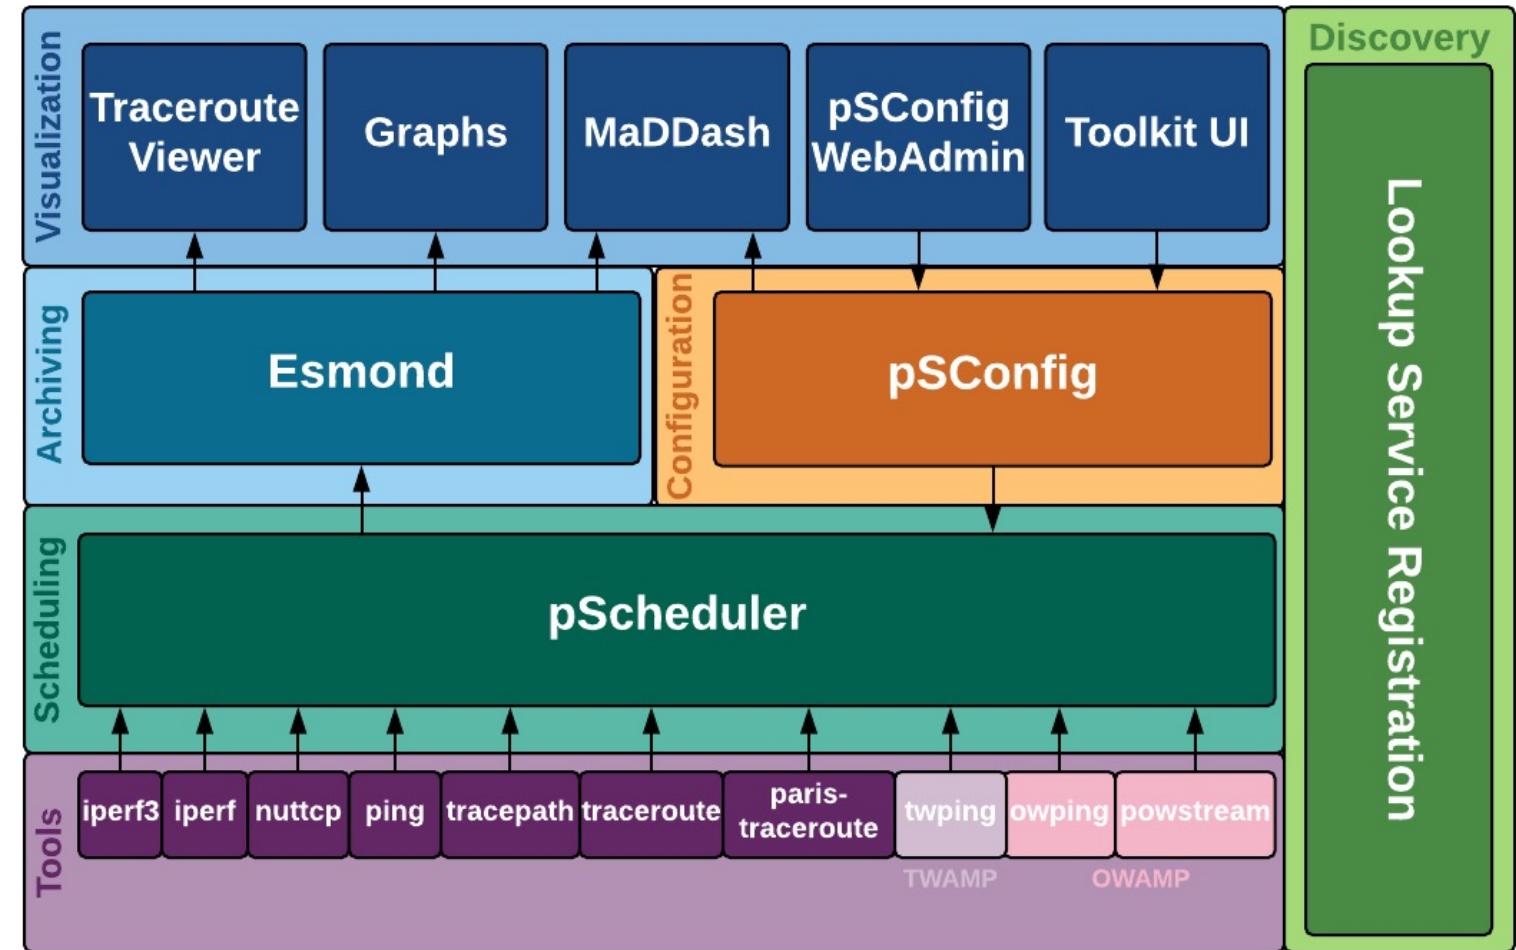
```
---  
apiVersion: k8s-netem.riasc.eu/v1  
kind: TrafficProfile  
metadata:  
  name: profile-delay-jitter  
spec:  
  podSelector:  
    matchLabels:  
      application: py_ctrl  
  
type: Builtin  
parameters:  
  netem:  
    delay: 0.2 # seconds  
    jitter: 0.05  
    loss_ratio: 1 # in [0, 1]  
    distribution: "normal"
```

```
egress:  
  - to:  
    - ipBlock:  
        cidr: 1.1.1.1/32  
    - podSelector:  
        matchLabels:  
          component: example  
  
  - to:  
    - ipBlock:  
        cidr: 192.168.0.44/32  
    - ports:  
      - port: 12000  
      protocol: UDP
```



2. Network Monitoring

- Work-in-Progress
- Based on perfSONAR
- Operator for configuring test schedules via Kubernetes manifests
- Testpoints run on nodes in the labs
- Central cloud is running web interfaces and measurement archive



2. Network Monitoring /2

```
---  
apiVersion: riasc.eu/v1  
kind: SchedulerConfig  
metadata:  
  name: test  
spec:  
  groups:  
    latency_group:  
      type: mesh  
      addresses:  
        - name: rpi-rwth-1  
        - name: ntnu-1  
  tests:  
    latency_test:  
      type: latencybg  
      spec:  
        source: "% address[0] %"  
        dest: "% address[1] %"  
        packet-interval: 0.1  
        packet-count: 600
```

```
archives:  
  esmond_archive:  
    archiver: esmond  
    data:  
      url: https://esmond.riasc-system  
      measurement-agent: "% scheduled_by_address %"  
schedules:  
  every_4_hours:  
    repeat: PT4H  
    slip: PT4H  
    sliprand: true  
tasks:  
  latency_task:  
    group: latency_group  
    test: latency_test  
    archives:  
      - esmond_archive
```



3. Time Synchronization

- Runs a dedicated Pod on nodes with
 - Chrony (NTP), GPSd and ptpt4linux containers
- Direct hardware access
 - Network cards (PTP), GPS Receivers, PPS Inputs
- Synchronization sources:
 - NTP: Network Time Protocol
 - GPS: Global Positioning System
 - PPS: Pulse per Second trigger
 - PTP: Precision Time Protocol
- Configured via Helm chart
- Can act as NTP server and client
 - Redistribute time within a lab...



Time Synchronization /3

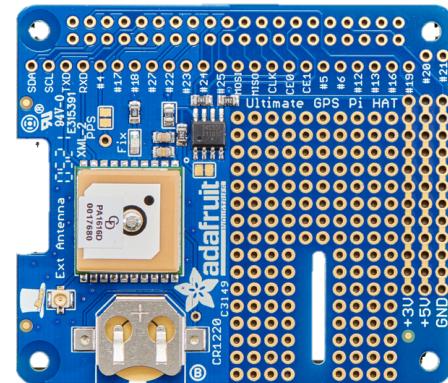
```
---  
apiVersion: riasc.eu/v1  
kind: TimeSyncConfig  
metadata:  
  name: rwth  
spec:  
  nodeSelector:  
    kubernetes.io/hostname: rpi-rwth-1  
  
gps:  
  enabled: true  
  device: ttyAMA0  
  
pps:  
  enabled: true  
  device: pps0  
  pin: 18
```

```
ntp:  
  server:  
    enabled: true  
    local: true  
    stratum: 1  
    allow:  
      - 134.130.169.0/25  
  
servers:  
- address: ntp1.rwth-aachen.de  
- address: ntp.acs-lab
```

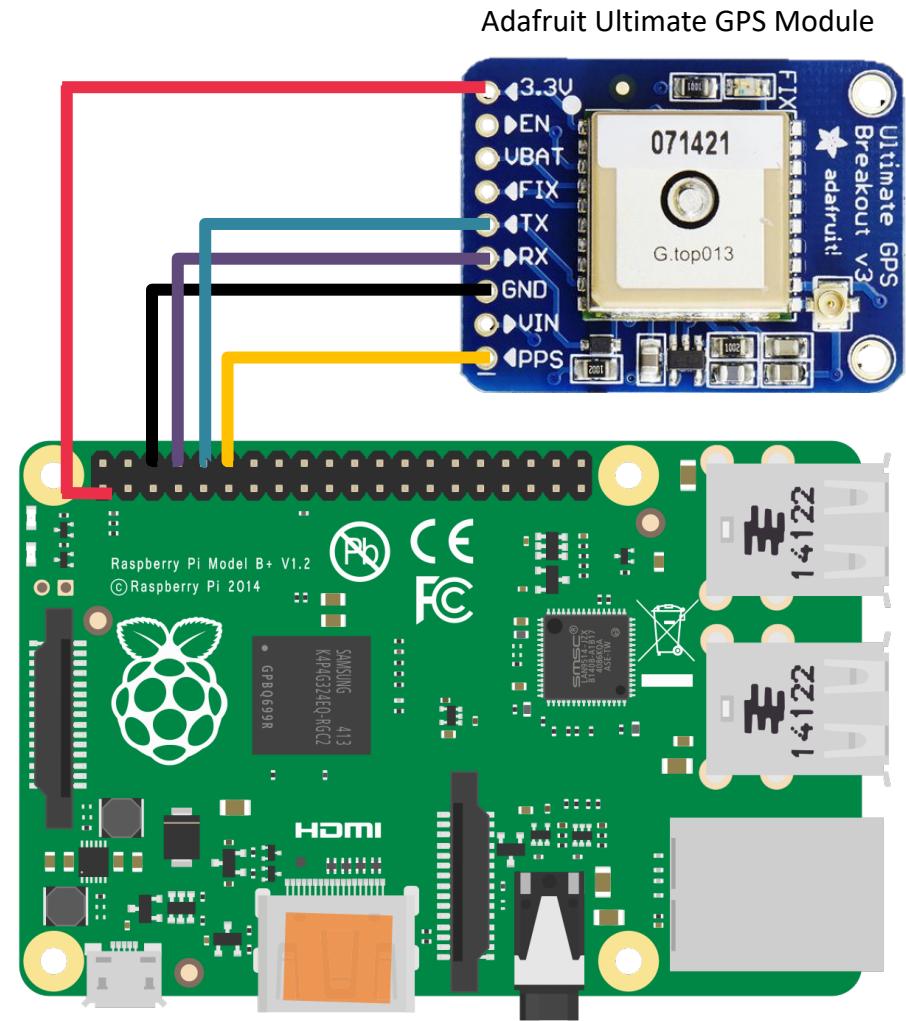


Time Synchronization /2

- Low Cost Solution for Raspberry Pi
- Adafruit Ultimate GPS Module v3 (~30 USD)
 - Or Adafruit Ultimate GPS Hat
- Provides Time-of-Day via UART port (NMEA)
- Provides Start-of-Second via Pulse-per-Second output (PPS)
- Raspberry Pis have no support for PTP (IEEE1588)



Adafruit Ultimate GPS Hat



Adafruit Ultimate GPS Module



Multi-tenancy

- Allows to restrict access of a user-group to a limited set of nodes
- Implemented as an operator and CRD
- Creates Kubernetes Namespace and Access control lists for us

```
---
```

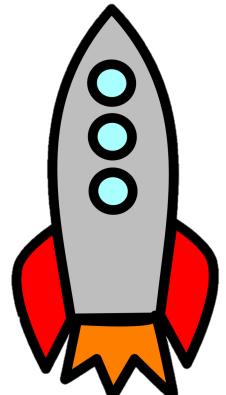
```
apiVersion: riasc.eu/v1
kind: Project
metadata:
  name: sim-hes-off
spec:
  users:
    - steffen-vogel
    - erick-alves

  nodes:
    - rpi-rwth-1
    - ntnu-1
```



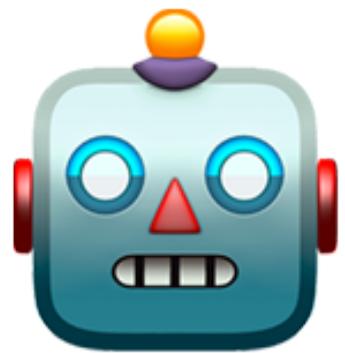
Getting Started

1. Provision your own node(s)
 - <https://riasc.eu/docs/setup/agent>
2. Request user account & project for Cluster access
 - Mail: svogel2@eonerc.rwth-aachen.de
3. Read Quickstart guide
 - <https://riasc.eu/docs/usage> (WIP)



Write your own Operator

- kopf: Kubernetes Operator Framework
 - Write operators easily in Python
 - Single file per Operator
- Recipe:
 1. Define ressource description (CRD)
 2. Implement operator logic
 3. Run / Deploy operator
 - Locally: “kopf run my_operator.py -verbose”
 - As a container within the cluster itself



Example: Device Configuration

CRD Definition

```
---  
apiVersion: apiextensions.k8s.io/v1  
kind: CustomResourceDefinition  
metadata:  
  name: chromq4qs.device.riasc.eu  
spec:  
  scope: Cluster  
  group: device.riasc.eu  
  names:  
    kind: Chroma4Q  
    plural: chroma4qs  
    singular: chroma4q  
  
versions:  
- name: v1  
  served: true  
  storage: true  
  schema:  
    openAPIV3Schema:  
      type: object  
      properties:  
        spec:  
          type: object  
          properties:  
            state:  
              type: string  
              default: disconnected  
              pattern: '^^(dis)?connected$'  
[ ... ]
```



Example: Device Configuration

Example Manifest



```
---  
apiVersion: device.riasc.eu/v1  
kind: Chroma4Q  
metadata:  
  name: rwth-chroma  
  namespace: sim-hes-off  
spec:  
  connection:  
    host: chroma.acs-lab  
    port: 2101  
    timeout: 20  
  
  state: disconnected  
  phases: [1, 2]
```

```
parameters:  
  maxCurrent: 16.0 # A  
  overcurrentDelay: 0.0 # s  
  maxPower: 3500 # VA  
  maxFrequency: 55.0 # Hz  
  maxVoltageAC: 235.0 # V  
  
setpoints:  
  voltageDC: [ 0.0, 0.0 ]  
  voltageRMS: [ 230.0, 240.0 ]  
  frequency: [ 50.0, 51.0 ]
```



Example: Device Configuration

Operator Logic: Creation

```
import kopf

@kopf.on.create('device.riasc.eu', 'v1', 'chroma4qs')
@kopf.on.resume('device.riasc.eu', 'v1', 'chroma4qs')
def create_or_resume(spec: kopf.Spec, memo: kopf.Memo):
    conn = spec.get('connection')
    params = spec.get('parameters')
    setp = spec.get('setpoints')
    if params is None or conn is None:
        raise kopf.PermanentError('incomplete settings')

    memo.amp = chroma.amp4Q(
        conn.get('host'),
        conn.get('port'),
        conn.get('timeout'), False)

    configure(memo.amp, params, setp)
```



Example: Device Configuration

Operator Logic: Update & Deletion

```
@kopf.on.update('device.riasc.eu', 'v1', 'chroma4qs')
def update(spec: kopf.Spec, memo: kopf.Memo):
    params = spec.get('parameters')
    setp = spec.get('setpoints')
    if params is None or setp is None:
        raise kopf.PermanentError('incomplete settings')
```

```
configure(memo.amp, params, setp)
```

```
@kopf.on.delete('device.riasc.eu', 'v1', 'chroma4qs')
def delete(memo: kopf.Memo):
    memo.amp.disconnect_DUT()
```



Example: Device Configuration

Operator Logic: Update & Deletion

```
@kopf.on.update('device.riasc.eu', 'v1', 'chroma4qs')
def update(spec: kopf.Spec, memo: kopf.Memo):
    params = spec.get('parameters')
    setp = spec.get('setpoints')
    if params is None or setp is None:
        raise kopf.PermanentError('incomplete settings')
```

```
configure(memo.amp, params, setp)
```

```
@kopf.on.delete('device.riasc.eu', 'v1', 'chroma4qs')
def delete(memo: kopf.Memo):
    memo.amp.disconnect_DUT()
```



Example: Device Configuration

Operator Logic: Monitoring

```
@kopf.timer('device.riasc.eu', 'v1', 'chroma4qs', interval=5)
def measurements(spec: kopf.Spec, memo: kopf.Memo):
    phases = spec.get('phases')

    return {
        'frequency': [memo.amp.meas_frequency(i)           for i in phases],
        'voltageAC': [memo.amp.meas_voltage_AC(i)         for i in phases],
        'currentDC': [memo.amp.meas_current_AC(i)          for i in phases],
        'powerReal': [memo.amp.meas_power_real(i)          for i in phases],
        'powerReactive':[memo.amp.meas_power_reactive(i)   for i in phases],
    }
```



Outlook

- More operators
 - VILLASnode
 - perfSonar
- Full experiment configuration with a single file
 - Helm?

Now: Feedback, Questions & Discussion



Links

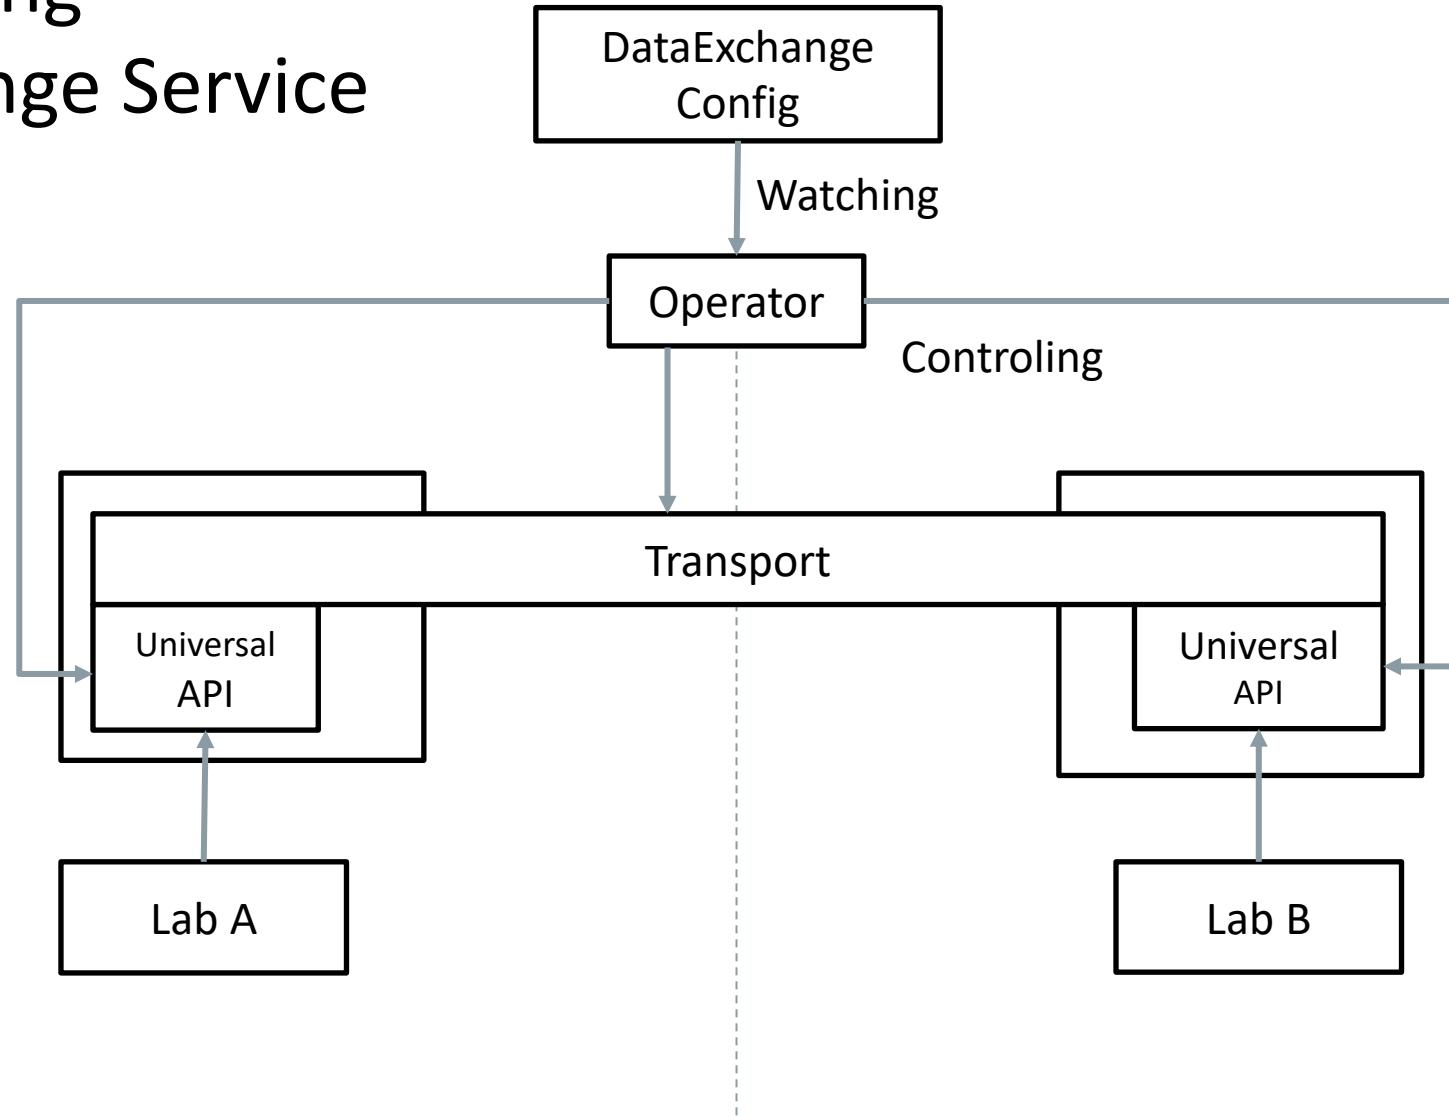


- <https://riasc.eu>
- <https://github.com/orgs/ERIGrid2/teams/riasc/repositories>

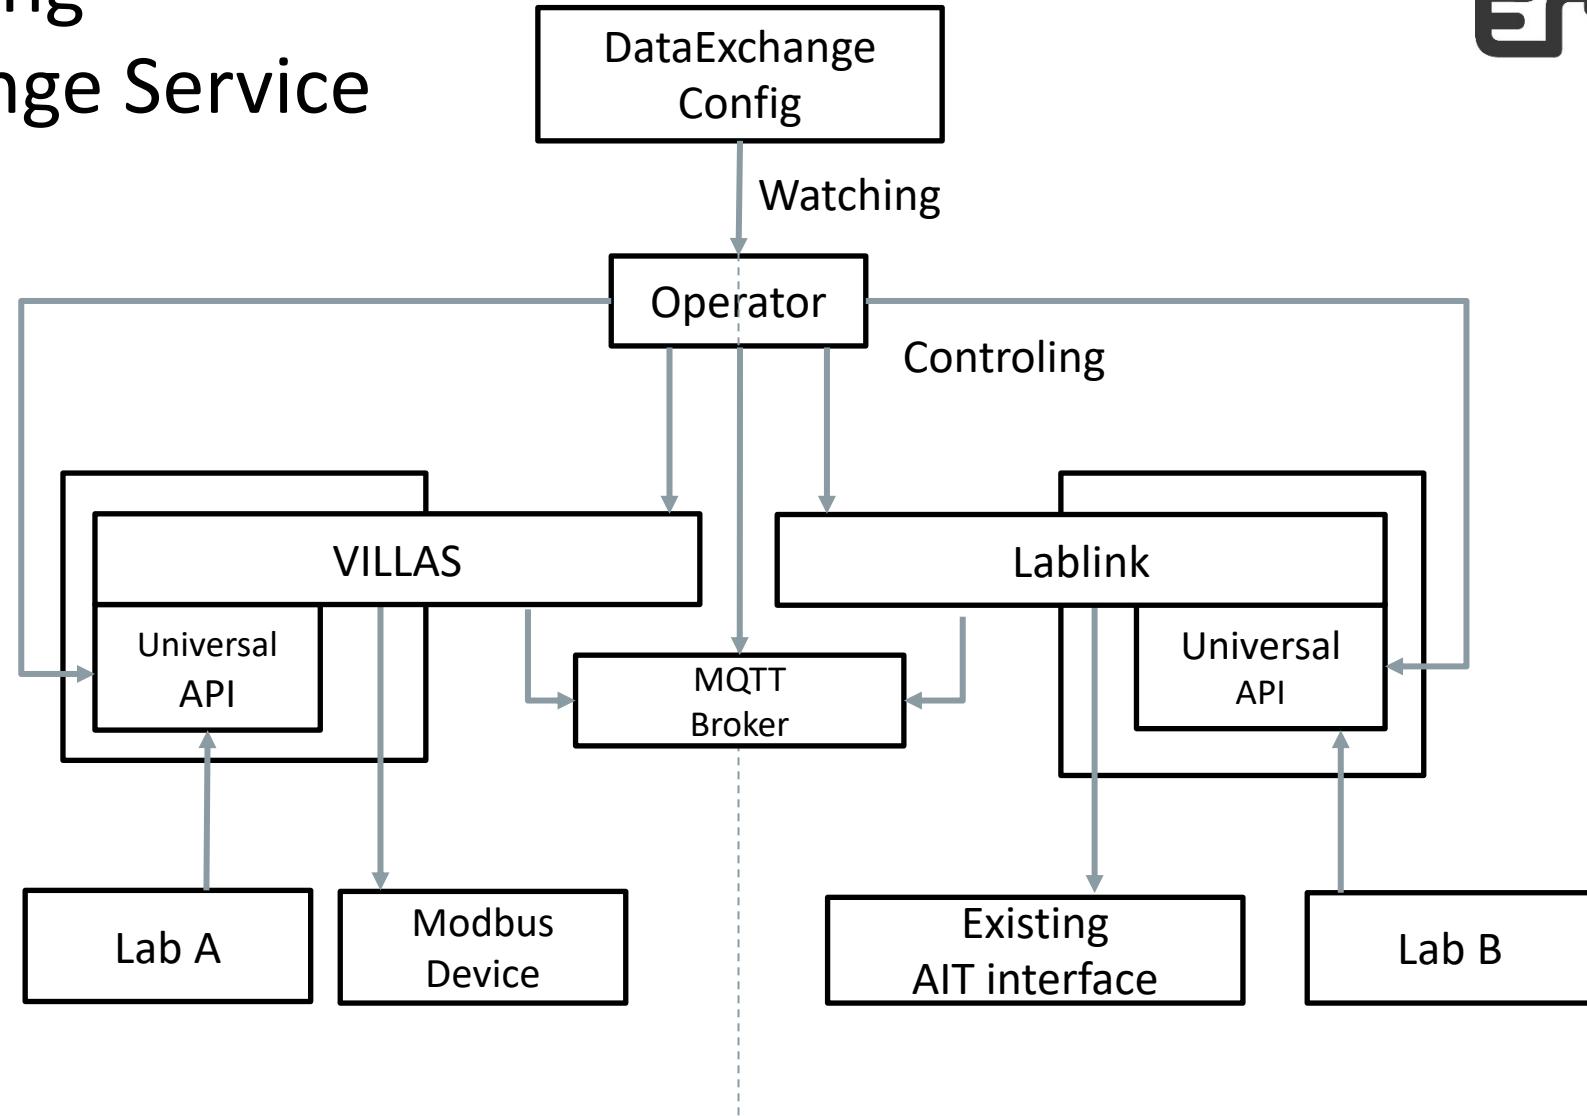
- <https://kubernetes.io>



Brainstorming Data Exchange Service



Brainstorming Data Exchange Service



Helm

